



Paired benchmark differences: DAGFormer versus baseline

● 75M ■ 150M ▲ 300M

(a) Accuracy endpoints



(b) Likelihood endpoints



95% document-bootstrap intervals; fixed checkpoints; positive favors DAGFormer.

* Strong label bias: all models are below the best constant-label score on these tasks.